

PRODUCT DISCOVERY ARTIFACT

Version 1.0 —

Product: Tattava Copilot

Company: Don Vanzara Showbiz LLP / TattvaCo

Discovery Focus: V1 — Story Development Copilot / Narrative Intelligence

Prepared by: Kaustubh Deshmukh — Product & AI Strategy Lead

Date: September 2026

Status: Discovery Baseline — User Validation

1. Executive Summary

Tattava Copilot is envisioned as an **AI-native film development and production copilot** that supports the journey from an initial creative idea toward a production-ready, produced and localized project.

The BRD establishes a long-term product journey:

Idea → Research → Story → Screenplay → Dialogue → Story World → Visual Development → Production Planning → Production → Post-Production → Final Content

The central architectural/product principle is that project intelligence should persist throughout this lifecycle rather than being recreated independently at every stage.

For V1, the immediate focus is:

V1 — Story Development Copilot / Narrative Intelligence

The proposed V1 transforms:

Initial creative input → structured Story Development Package

through a workflow involving:

Input → Story Brain → Research → Character Intelligence → Narrative Reasoning → Story Development → Format/Genre → Template → Writing → Screenplay → Dialogue → Story Doctor → Canon/Continuity → Evaluation

However, the BRD is a **business/product direction document**, not evidence that every proposed capability has been validated with users.

Therefore, the purpose of this Product Discovery Artifact is to answer:

Who has the problem?

What exactly is the problem?

How does it occur in the real workflow?

How painful and frequent is it?

How do users solve it today?

Where can AI create differentiated value?

What will users trust AI to do?

What is the smallest useful V1?

2. Discovery Principle

The product development process should follow:

```
BRD
↓
Discovery
↓
Evidence
↓
Validated Problem
↓
JTBD
↓
Opportunity
```

```
↓  
Hypothesis  
↓  
Experiment  
↓  
Validation  
↓  
Product Decision  
↓  
Product Requirement  
↓  
Feature  
↓  
Acceptance Criteria  
↓  
Metric
```

Not:

```
BRD  
↓  
Features  
↓  
Engineering
```

This distinction is critical for an AI-native product because the existence of an AI capability does not automatically mean there is a meaningful user problem for it to solve.

3. Source Baseline — What the BRD Establishes

The BRD establishes the following strategic foundation.

3.1 Product vision

Tattava is intended to become:

An AI-native operating system for film development and production that transforms ideas into structured, production-ready projects while keeping creative teams in control.

3.2 Business context

Tattava has two connected but distinct contexts:

Venture 1 — AI-Native Content / Tattava

Beginning with:

- Story development
- Visual development
- AI filmmaking
- Localization

Venture 2 — Conventional Film Production

Traditional/live-action filmmaking that can still benefit from:

- intelligence
- planning
- production management
- connected project information

The BRD specifically states that the production layer should be technology-agnostic so AI-generated and conventional films can eventually enter the same production operating system.

4. V1 Strategic Scope

The BRD identifies V1 as:

| Story Development Copilot — Narrative Intelligence

Its objective is to transform an initial creative input into a structured story-development package supported by persistent project intelligence.

The BRD defines the Story Brain as the persistent project-intelligence core containing information such as:

- premise
- characters
- world
- decisions
- research

and making this context available downstream.

5. Product Discovery Status

Discovery Area	Current Status	Confidence
Strategic product vision	Defined in BRD	High
V1 direction	Defined	High
Primary user categories	Defined	Medium
Current workflow	Mostly inferred	Low
Pain severity	Not quantified	Low
Pain frequency	Not quantified	Low
Story Brain value	Strategic hypothesis	Medium
Research value	Strategic requirement	Medium
Continuity value	Strategic requirement	Medium
AI trust	Not validated	Low
Existing AI usage	Unknown	Low
Switching behavior	Unknown	Low
ROI	Unknown	Low
Exact V1 MVP	Not validated	Low

Discovery conclusion

Tattava has a strong strategic direction but does not yet have sufficient empirical evidence to treat the complete V1 scope as validated.

6. Opportunity Brief

6.1 Opportunity statement

There is an opportunity to create a **persistent AI-native project intelligence environment for entertainment development** that reduces fragmentation between:

- research
- story development
- character development
- writing
- screenplay development
- dialogue
- revision
- continuity
- evaluation
- project packaging

The BRD identifies fragmentation across people, documents, tools and stages as a central problem.

7. Opportunity Reframing

The opportunity should not be framed simply as:

| **"Build an AI script-writing tool."**

A stronger discovery framing is:

| **"How might we help entertainment teams continuously develop a project without repeatedly losing context, reconstructing information, duplicating work or manually checking continuity?"**

This moves the product from:

AI generation

toward:

AI-native project intelligence.

8. Potential Opportunity Areas

Opportunity 1 — Context fragmentation

Project knowledge may be distributed across:

- documents
- scripts
- conversations
- notes
- research
- meetings
- people
- AI chats

Hypothesis

A persistent project intelligence layer can reduce repeated context-setting.

Opportunity 2 — Repetitive work

Creative teams may repeatedly:

- research
- summarize
- restructure
- rewrite

- convert
- document
- compare
- consolidate

Hypothesis

AI can automate parts of this transformation while preserving human control.

Opportunity 3 — Research traceability

Creative development often requires:

- historical research
- cultural research
- subject-matter research
- genre research
- market/context research
- audience research

The BRD requires research to capture source, reference, finding and relevance.

Hypothesis

Research becomes more useful when evidence is directly connected to project decisions.

Opportunity 4 — Continuity

The product is intended to maintain canon across:

- characters
- world
- timeline
- story logic

and identify contradictions.

Hypothesis

AI-assisted continuity can reduce manual review effort.

Opportunity 5 — Handoff friction

The project moves through:

```
graph TD; Research --> Story; Story --> Screenplay; Screenplay --> Dialogue; Dialogue --> VisualDevelopment[Visual Development]; VisualDevelopment --> Production;
```

Research
↓
Story
↓
Screenplay
↓
Dialogue
↓
Visual Development
↓
Production

The strategic opportunity is to preserve context across those transitions.

9. Problem Definition

9.1 Primary problem statement

Entertainment development teams need to continuously transform and refine a project across research, story, characters, writing, screenplay, dialogue and review, but the knowledge and decisions created during this process can become fragmented across people, tools and artifacts, creating repetitive work, context switching, continuity challenges and difficulty maintaining a coherent project state.

This is consistent with the BRD's problem statement around fragmentation across people, documents, tools and stages.

10. Problem Decomposition

Problem	Description	Validation
Research fragmentation	Research may be separated from creative decisions	Required
Project-context fragmentation	Users may repeatedly restate project context	Required
Repetitive transformation	Same knowledge may be converted across artifacts	Required
Continuity risk	Established facts may become inconsistent	Required
Collaboration friction	Multiple stakeholders contribute to development	Required
AI trust	Users need confidence before accepting AI output	Required
Version/revision complexity	Creative projects evolve through many drafts	Required

11. Target User Model

The BRD identifies primary users as:

- Directors
- Writers
- Story developers
- Creative producers
- Showrunners
- Content development teams

Secondary users include:

- Art directors
- Production designers
- Casting teams

- Production managers
 - Line producers
 - Producers
 - Business/operations teams.
-

12. Discovery User Segmentation

P0 — Writers / Screenwriters

Why:

- closest to writing
 - research
 - story development
 - revision
 - AI experimentation
-

P0 — Story Developers / Story Editors

Why:

- narrative structure
 - character development
 - story logic
 - continuity
 - evaluation
-

P0 — Creative Producers

Why:

- development coordination

- creative decision-making
 - progression
 - packaging
-

P0 — Directors

Why:

- creative direction
 - narrative interpretation
 - approval
 - downstream visual intent
-

P0 — Showrunners / Content Leads

Why:

- story/world decisions
 - episode-level development
 - creative oversight
-

P1 — Producers

Important for progression and future production readiness.

P2 — Production / Art / Specialist Teams

Important primarily for future versions and downstream intelligence.

13. User Role Hypothesis

Role	Likely responsibility	Confidence
Writer	Story, screenplay, dialogue	Medium

Role	Likely responsibility	Confidence
Story Developer	Narrative development	Medium
Story Editor	Structure, review, continuity	Medium
Director	Creative direction	Medium
Creative Producer	Coordination and progression	Medium
Showrunner	Story/world/episode direction	Medium
Producer	Project progression	Medium

These are **working hypotheses**, not validated organizational mappings.

14. Discovery Research Questions — Preliminary Answers

USER

1. Who actually performs story development?

Story development appears to be a **multi-role activity** involving:

- writers
- story developers
- story editors
- directors
- creative producers
- showrunners
- content-development teams

The exact division of responsibility is unknown.

2. Who owns creative decisions?

The creative team owns final decisions.

The BRD explicitly establishes human-in-the-loop interaction and says AI should not independently make irreversible creative decisions.

Exact approval ownership by stage remains unknown.

3. Who performs research?

Likely:

- writers
- story developers
- creative producers
- researchers/knowledge-management specialists where applicable

The BRD does not establish a dedicated research role.

4. Who reviews drafts?

Likely:

- writers
- story developers/editors
- directors
- creative producers
- showrunners
- other creative stakeholders

Exact review hierarchy requires research.

5. Who approves progression?

Human creative stakeholders.

The AI should recommend/assist, but the human should approve.

6. Who collaborates with writers?

Potentially:

- directors
- story developers
- story editors
- producers
- creative producers
- showrunners
- researchers
- subject-matter experts

The actual collaboration model must be observed.

15. WORKFLOW

7. What happens from idea to first structured story?

The **proposed Tattava workflow** is:

```
graph TD; A[Idea / Input] --> B[Intent Understanding]; B --> C[Story Brain]; C --> D[Research]; D --> E[Character Intelligence]; E --> F[Narrative Reasoning]; F --> G[Story Development]; G --> H[Format / Genre]; H --> I[Template]; I --> J[Writing]; J --> K[Screenplay]; K --> L[ ];
```

```
Dialogue
  ↓
Story Doctor
  ↓
Canon / Continuity
  ↓
Evaluation
  ↓
Final Story Development Package
```

This is what Tattava proposes, **not a validated representation of how users currently work.**

8. What happens after the first story version?

Working hypothesis:

```
Draft
  ↓
Review
  ↓
Feedback
  ↓
Revision
  ↓
Continuity
  ↓
Evaluation
  ↓
Approval
  ↓
Next version
```

The BRD explicitly supports review, editing, rewriting, refinement, approval and regeneration.

Actual iteration frequency is unknown.

9. Where are decisions recorded?

The future Tattava model proposes that important decisions become part of persistent project intelligence.

Current-state discovery must identify whether decisions are currently stored in:

- documents
 - notes
 - scripts
 - meeting records
 - messaging
 - emails
 - AI chats
 - people's memory
-

10. How do teams collaborate?

Likely through a mixture of:

- meetings
- documents
- reviews
- feedback
- conversations
- creative approvals

But the current collaboration model is not documented in the BRD.

11. Where are drafts stored?

Unknown.

Discovery must identify the actual systems and the current source of truth.

12. How are revisions managed?

The product requires version control, but current-state revision management is unknown.

Discovery should investigate:

- version naming
- feedback
- comparison
- approval
- rollback
- draft ownership
- canonical version

The BRD identifies version control as a non-functional requirement.

16. PAIN

13. Where is the most time spent?

The BRD identifies:

- research
- story development
- writing
- repetitive work
- structuring
- documenting
- transforming material

as areas where friction may exist.

But we **do not yet know which consumes the most time.**

This must be measured.

14. Where is work duplicated?

Potential chain:

```
graph TD; Research --> Notes; Notes --> Story; Story --> Synopsis; Synopsis --> Characters; Characters --> Treatment; Treatment --> Screenplay; Screenplay --> Dialogue; Dialogue --> Presentation;
```

Research
↓
Notes
↓
Story
↓
Synopsis
↓
Characters
↓
Treatment
↓
Screenplay
↓
Dialogue
↓
Presentation

This is a hypothesis.

15. Where does information get lost?

Potentially across:

- research → story
- story → screenplay
- character → dialogue
- meeting → documentation
- writer → director
- development → production

The BRD's central problem hypothesis is precisely that decisions and context can become disconnected as projects move between people, tools and stages.

16. Which tasks are most frustrating?

Unknown.

We should not invent the answer.

Candidates to investigate:

- repetitive research
 - revisions
 - continuity
 - feedback incorporation
 - finding previous decisions
 - documentation
 - formatting
 - collaboration
-

17. Which tasks create delays?

Unknown.

Candidate areas:

- research
- feedback
- revision
- creative alignment
- approvals
- information retrieval
- continuity checking

Research should identify the actual bottleneck.

17. AI

18. Where are users already using AI?

Discovery must ask users to show actual AI usage.

Potential categories:

- brainstorming
 - research
 - rewriting
 - summarization
 - character development
 - screenplay
 - dialogue
 - image generation
 - reference generation
 - analysis
-

19. Which AI tasks do users trust?

Working hypothesis:

Higher trust may exist for:

- brainstorming
- summarization
- rewriting
- alternative generation
- research assistance

Lower trust may exist for:

- factual claims
- final creative decisions
- culturally sensitive content

- continuity-critical output

Must be validated.

20. Which AI tasks do they not trust?

Potential candidates:

- hallucinated research
- generic storytelling
- inconsistent characters
- culturally inaccurate content
- unsupported factual claims
- autonomous creative decisions

The BRD identifies creative, continuity, adoption and data/IP risks.

21. What causes users to verify AI output?

Likely triggers:

Factual uncertainty

"Is this actually true?"

Source uncertainty

"Where did this come from?"

Creative mismatch

"Does this fit my story?"

Continuity uncertainty

"Does this contradict something we already established?"

High-consequence decision

"Can this safely enter the final project?"

22. What would make AI unacceptable?

Potential rejection triggers:

- hallucination
- generic output
- loss of creative control
- IP concerns
- privacy/security concerns
- cultural inaccuracies
- poor continuity
- lack of source traceability
- inability to edit/reject
- unpredictable behavior

These remain hypotheses.

18. PRODUCT

23. Which workflow should Tattava own?

Preliminary answer

The strongest strategic candidate is:

| **The connected project-development intelligence layer.**

Not simply:

| "AI screenplay generation."

Potential ownership:

```
graph TD; Research --> PI[Project Intelligence]; PI --> SD[Story Development]; SD --> CI[Character Intelligence]; CI --> Writing; Writing --> Screenplay; Screenplay --> Dialogue; Dialogue --> CC[Canon / Continuity]; CC --> Evaluation;
```

The BRD explicitly positions the Story Brain and Project Intelligence Layer as foundational across later product stages.

24. Which existing tools should Tattava integrate with?

Do not decide integrations before understanding:

- actual tools
 - frequency of use
 - source-of-truth systems
 - switching costs
 - export/import needs
 - collaboration patterns
-

25. What would cause someone to change their workflow?

Potential triggers:

- significant time savings

- less repetitive work
- better project memory
- stronger continuity
- better research traceability
- faster story development
- fewer revisions
- easier collaboration
- better structured outputs

The BRD identifies acceleration, reduced development friction and improved continuity as business objectives.

But whether these are sufficient to cause switching is unknown.

19. What Is the Smallest Useful Product?

This is currently the **most important unresolved question**.

The BRD describes a broad V1 containing:

```
Input
→ Story Brain
→ Research
→ Character Intelligence
→ Narrative Reasoning
→ Story Development
→ Format/Genre
→ Template
→ Writing
→ Screenplay
→ Dialogue
→ Story Doctor
→ Canon/Continuity
→ Evaluation
→ Presentation
```


Discovery should determine whether all of these are required for the first release.

20. Current MVP Hypothesis

The current discovery hypothesis is:

The smallest potentially meaningful Tattava product is Project Intelligence + Research + Context-Aware Story Development + Canon/Continuity, wrapped in a simple human-controlled creative workflow.

Conceptually:



This is **not yet the approved MVP**.

It is the hypothesis that discovery should test.

21. Jobs To Be Done

JTBD-01 — Develop an Idea

Situation

I have an early creative idea.

Job

I want to explore and structure it so I can understand whether it has enough narrative potential to develop further.

Desired outcome

- premise
 - conflict
 - characters
 - themes
 - narrative possibilities
 - open questions
-

22. JTBD-02 — Conduct Meaningful Research

Job

When my story depends on unfamiliar historical, cultural, technical or contextual information, I want reliable evidence connected to my project so that I can make informed creative decisions.

Desired outcome

Relevant + credible + traceable research.

The BRD explicitly requires research to capture source, reference, key finding and relevance.

23. JTBD-03 — Maintain Project Context

When my project evolves through multiple drafts and creative decisions, I want the system to remember approved project context so that I don't have to repeatedly explain the same things.

Desired outcome

Persistent project memory.

24. JTBD-04 — Develop Characters

When developing a character, I want their motivations, relationships and voice to remain coherent throughout the project.

Desired outcome

Consistent character intelligence.

25. JTBD-05 — Develop Story Structure

When shaping a story, I want to explore narrative structures and identify logical weaknesses before investing heavily in screenplay development.

Desired outcome

Better-informed narrative decisions.

26. JTBD-06 — Write With Context

When transforming approved story material into screenplay or dialogue, I want AI assistance grounded in my project's established context so that I spend less time correcting generic or inconsistent output.

27. JTBD-07 — Review and Diagnose

When I have a draft, I want to identify structural, pacing, dialogue and continuity problems so that I can focus my creative effort on the issues that matter most.

The BRD defines Story Doctor and Evaluation for this purpose.

28. JTBD-08 — Prepare a Development Package

When a project is sufficiently developed, I want to consolidate its creative and research outputs into a structured package so that it can be reviewed and progressed.

The proposed V1 final output contains concept, research, logline, synopsis, genre/format, structure, characters, screenplay, dialogue, research, evaluation and project metadata.


29. Opportunity Tree


North Star

Reduce the time and operational complexity required to move an entertainment idea from concept toward production-ready execution while preserving human creative control.

The BRD explicitly defines this as the North Star.

Opportunity structure





30. Assumption Map

Assumption	Impact	Uncertainty	Validation
Context fragmentation is a major pain	High	High	User interviews
Story Brain solves meaningful pain	High	High	Prototype
Users want AI integrated into workflow	High	High	Current behavior study
Traceable research increases trust	High	High	Research experiment
Continuity checking saves time	High	High	Task test
Integrated workflow beats point tools	High	High	End-to-end prototype
Users will change workflows	High	High	Switching interviews
V1 can create value independently	High	Medium	V1 prototype

31. Hypothesis Register

H1 — Persistent Project Intelligence

Hypothesis

Users repeatedly lose or restate project context.

Expected impact

Persistent memory reduces:

- repeated prompting
- context switching
- manual reconstruction

Test

Give users a project-memory prototype.

Measure

- task completion time
 - repeated context input
 - corrections
 - perceived usefulness
-

32. H2 — Traceable Research

Hypothesis

Research with visible sources and project relevance is more trusted than generic AI research.

Test

Compare:

A: Answer only

B: Answer + sources

C: Answer + sources + project relevance

Measure

- trust
- verification

- usefulness
 - willingness to use
-

33. H3 — Context-Aware Generation

Hypothesis

Project-grounded AI produces more usable output than generic prompting.

Measure

- acceptance rate
 - edits
 - corrections
 - regeneration
 - time-to-usable-output
-

34. H4 — Continuity Intelligence

Hypothesis

AI-assisted continuity checking reduces manual review effort.

Test

Give users a project containing deliberate inconsistencies.

Measure

- issues detected
 - false positives
 - missed issues
 - time required
-

35. H5 — Human-Controlled AI

Hypothesis

Creative professionals prefer:

```
AI proposes
↓
Human evaluates
↓
Human modifies
↓
Human approves
↓
Project updates
```

rather than:

```
AI decides
↓
Project changes automatically
```

This is directly aligned with the BRD's human-in-the-loop principle.

36. H6 — Connected Workflow

Hypothesis

The combination of:

Research + Project Intelligence + Story Development + Writing + Continuity

creates more value than isolated AI tools.

This is arguably the most strategically important V1 hypothesis.

37. Discovery Research Plan

Objective

Validate:

1. actual user workflow
 2. actual pain
 3. actual frequency
 4. actual severity
 5. current tools
 6. current AI usage
 7. trust boundaries
 8. workflow-switching barriers
 9. value of persistent project intelligence
 10. smallest useful V1
-

38. Recommended Research Participants

Segment	Initial sample
Writers / Screenwriters	4–6
Story Developers / Editors	3–4
Creative Producers	2–3
Directors / Showrunners	2–3
Initial total	11–16

Expand if significant new patterns continue emerging.

39. Interview Method

Do **not** start by explaining Tattava.

Do **not** ask:

| "Would you use an AI Story Brain?"

Do **not** ask:

| "Would you pay for Tattava?"

Instead, reconstruct a **recent real project**.

40. Interview Script

Step 1 — Recent project

| "Tell me about the most recent project you worked on."

Step 2 — Workflow

| "Walk me through what happened from the original idea to the first serious draft."

Step 3 — People

| "Who was involved at each stage?"

Step 4 — Tools

| "Where did you keep the information?"

Step 5 — Decisions

| "What were the important decisions and where were they recorded?"

Step 6 — Pain

| "Which part took longer than it should have?"

Step 7 — Evidence

| "Can you show me how you handled that?"

Step 8 — Failure

| "Tell me about a time something went wrong."

Step 9 — AI

| "Did you use AI anywhere? Show me how."

Step 10 — Trust

| "What did you trust AI to do?"

Step 11 — Verification

| "What did you verify and why?"

Step 12 — Ideal state

| "If you could remove one frustrating part of this process, what would it be?"

41. Evidence Capture Framework

For every interview:

Field	Capture
Participant	Name/role
Project	Recent project
Stage	Workflow stage
Task	What they were trying to accomplish
Current tool	Tool/workaround
Behavior	What they actually did
Problem	Pain point
Frequency	How often
Severity	Impact
Time	Time spent

Field	Capture
Quote	Direct evidence
AI usage	Existing AI behavior
Trust barrier	What they verify
Opportunity	Potential product opportunity

42. Evidence Hierarchy

L0 — Opinion

| "I think this is a problem."

Low evidence.

L1 — Stated problem

| "I frequently lose research."

Useful but still self-reported.

L2 — Observed behavior

User demonstrates their workaround.

Stronger.

L3 — Repeated evidence

Multiple users show the same behavior.

High confidence.

L4 — Prototype evidence

Users successfully complete a task with the proposed solution.

Very strong.

L5 — Behavioral evidence

Real usage demonstrates measurable improvement.

Strongest.

43. Discovery Experiments

EXP-01 — Current Workflow Mapping

Goal

Understand actual idea-to-story workflow.

Method

Contextual interviews + artifact walkthrough.

Output

Current-state journey map.

EXP-02 — Story Brain Concept Test

Goal

Test whether persistent project memory creates value.

Prototype

Lightweight project-memory system.

Measure

- repeated context
- task time

- usefulness
 - corrections
-

EXP-03 — Research Agent Test

Goal

Test whether traceability improves trust.

Compare

Answer only
vs
Answer + Source
vs
Answer + Source + Project relevance

EXP-04 — Continuity Test

Goal

Determine whether AI can reduce manual continuity checking.

Method


Introduce deliberate contradictions.

Measure

- detection
 - false positives
 - time
 - usefulness
-

EXP-05 — End-to-End Workflow Test

Flow



Question

Does the connected workflow create substantially more value than using separate AI tools?

44. Current Workflow vs Tattava Hypothesis

Dimension	Current workflow hypothesis	Tattava hypothesis
Context	Repeatedly supplied	Persistent
Research	Separate	Project-connected
Sources	Potentially disconnected	Traceable
Characters	Documents/notes	Structured intelligence
Story	Multiple artifacts	Connected project state
Writing	Repeated prompting	Context-aware
Continuity	Manual/ad hoc	AI-assisted
Revision	Manual	Project-aware
Decisions	Distributed	Project intelligence
Output	Multiple artifacts	Structured package

These are **hypotheses, not claims of superiority.**

45. AI-Native Product Test

A capability should not be called AI-native merely because an LLM is involved.

For every feature, ask:

1. Intelligence

What does the system understand?

2. Context

What project information does it use?

3. Reasoning

What does it infer?

4. Action

What does it produce or execute?

5. Memory

What persists?

6. Feedback

How does human feedback change project intelligence?

7. Evaluation

How does the system determine whether the output is useful?

8. Control

Where does the human approve, reject or modify?

This aligns with the BRD's progression from:

| AI assistance → AI generation → AI orchestration.

46. Product Differentiation Hypothesis

Tattava should not primarily differentiate through:

| Better prompts

or:

| More AI agents

or:

| More content generation.

The stronger potential differentiation is:

1. Project memory

The system understands the project.

2. Project grounding

AI outputs are grounded in the project.

3. Research traceability

Evidence is connected to sources.

4. Creative continuity

Research → story → screenplay → dialogue remain connected.

5. Canon


The system understands established facts.

6. Human control

AI assists without replacing creative authority.

47. V1 Value Hypothesis

The current candidate value chain is:



This should be validated before becoming the final V1 architecture.

48. V1 Final Output

The BRD defines the V1 output as a structured Story Development Package containing:

- Project concept
- Research report
- Logline
- Synopsis
- Genre and format
- Story structure
- Character overview
- Screenplay
- Dialogue

- Supporting research
- Evaluation summary
- Project metadata / Story Brain state.

The discovery question is:

Which of these outputs do users actually need for V1, and which can wait?

49. North Star Metric

The BRD defines the North Star as:

Reduce the time and operational complexity required to move an entertainment idea from concept to production-ready execution, while preserving human creative control.

For V1, the practical proxy should be:

Time to Human-Approved Story Development Package

Measured from:

Initial accepted project input

to:

Human-approved structured development package.

50. Supporting Metrics

Efficiency

- Research turnaround
- Time to first usable story
- Time spent on repetitive work
- Number of iterations

- Time spent reconstructing context

Quality

- Human acceptance rate
- Correction rate
- Continuity issues
- Research verification rate

Trust

- AI rejection rate
- Verification rate
- Approval rate
- Source interaction

Adoption

- Projects created
- Active projects
- Repeat usage
- Projects reaching package completion

The BRD already identifies time to first structured story package, research turnaround, development iterations, creative adoption and time saved on repetitive work as V1 measures.

51. Discovery Decision Gates

Gate 1 — Problem Validation

Question:

| Does the problem actually exist and matter?

Evidence required:

- repeated user evidence
 - observed behavior
 - severity
 - current workaround
-

Gate 2 — Opportunity Validation

Question:

| Is there a valuable opportunity worth solving?

Evidence:

- meaningful pain
 - sufficient frequency
 - measurable impact
 - inadequate current solution
-

Gate 3 — Solution Validation

Question:

| Does Tattava's proposed workflow actually improve the situation?

Evidence:

- prototype usability
 - task completion
 - time improvement
 - user acceptance
 - trust
-

Gate 4 — MVP Definition

Question:

What is the smallest product that creates meaningful value?

Evidence:

- validated P0 jobs
- validated workflow
- validated AI behavior
- validated user acceptance

Gate 5 — Engineering Readiness

Only after validation:

```
Problem
↓
JTBD
↓
Evidence
↓
Product decision
↓
Requirement
↓
Feature
↓
Acceptance criteria
↓
Engineering
```

52. What We Know vs What We Don't Know

 **We know**

From the BRD:

- Tattava's strategic vision
 - V1 direction
 - target user categories
 - human-in-the-loop principle
 - Story Brain concept
 - research traceability requirement
 - continuity requirement
 - proposed V1 workflow
 - proposed V1 outputs
 - North Star direction
-

Working hypotheses

- Project context fragmentation is significant.
 - Users repeatedly recreate context.
 - Research is a meaningful pain.
 - Continuity is a meaningful pain.
 - Connected AI workflow is more valuable than isolated generation.
 - Persistent project intelligence will create differentiated value.
 - Users want AI embedded in their creative workflow.
-

Unknown

- Exact current workflow
- Exact pain severity
- Largest time sink
- Current tools
- Current AI usage

- Trust boundaries
 - Switching triggers
 - Willingness to adopt
 - ROI
 - Exact P0 features
 - Smallest useful V1
-

53. What Discovery Must NOT Assume

Do not assume:

1. Every BRD capability belongs in MVP.
 2. Writers necessarily have the strongest pain.
 3. Story Brain is automatically the solution.
 4. Research is the largest bottleneck.
 5. Users will switch from ChatGPT.
 6. Users want maximum AI autonomy.
 7. AI screenplay generation is the primary value.
 8. Continuity is the most painful problem.
 9. Users will change their existing workflow.
 10. A large number of AI agents means a better product.
-

54. Preliminary V1 Recommendation

Based on the current BRD and without pretending that user validation has already happened:

Candidate V1 Core

Persistent Project Intelligence + Research + Context-Aware Story Development + Canon/Continuity + Human Approval

Supporting capabilities

```
graph TD; Input --> PI[Project Intelligence]; PI --> Research; Research --> SD[Story Development]; SD --> CI[Character Intelligence]; CI --> Writing; Writing --> Continuity; Continuity --> Evaluation; Evaluation --> SO[Structured Output];
```

The rest of the BRD's capabilities can remain part of the larger roadmap rather than automatically becoming V1 engineering scope.

The BRD itself emphasizes incremental development and states that the product should be developed stage by stage rather than attempting the complete production operating system in the first release.

55. Discovery-to-PRD Traceability

Every future product requirement should be traceable:

```
graph TD; USER --> JOB; JOB --> PROBLEM; PROBLEM --> EVIDENCE; EVIDENCE --> OPPORTUNITY;
```

↓
HYPOTHESIS
↓
EXPERIMENT
↓
VALIDATION
↓
PRODUCT DECISION
↓
REQUIREMENT
↓
FEATURE
↓
USER STORY
↓
ACCEPTANCE CRITERIA
↓
METRIC

Example:

Writer
↓
Needs reliable research
↓
Research is fragmented
↓
5/6 interviewed writers show same behavior
↓
Traceable research opportunity
↓
Sources increase trust
↓
Research prototype
↓
Users verify less / complete task faster
↓
Build Research Intelligence
↓
FR-RESEARCH-01
↓
Research Agent
↓
Acceptance criteria
↓
Research turnaround + trust

This is the level of traceability we should maintain throughout Tattava's NPD process.

56. Management Decisions Required

Management should align on five decisions before final V1 Product Definition.

Decision 1

Approve the discovery problem space

Fragmentation across research, story development, writing, revision and continuity.

Decision 2

Approve P0 discovery users

- Writers
- Story Developers / Editors
- Creative Producers
- Directors / Showrunners

Decision 3

Approve discovery methodology

Recent-project interviews + artifact walkthroughs + prototype validation.

Decision 4

Approve central product hypothesis

Persistent, project-grounded intelligence can materially improve creative development workflow.

Decision 5

Approve provisional V1 scope

The V1 scope remains provisional until user evidence determines the smallest validated product.

57. Final Discovery Thesis


If Tattava can maintain a persistent, trustworthy representation of a creative project and use that context to connect research, story development, writing and continuity assistance, then creative teams may be able to move from idea to a structured, human-approved story development package with less repetitive work and lower context-switching cost than their current workflow.


This is the **central hypothesis to prove**.

It is **not yet a validated product claim**.

58. Final Product Discovery Position

The current state of Tattava can therefore be represented as:





That gives Tattava a proper **AI-native Product Discovery foundation** before moving into Product Definition and Engineering.

The next artifact should logically be:

| **V1 Product Definition & Product Blueprint v1.0**

and it should be created **only from the validated findings from this Discovery Artifact**, with every P0 capability traceable back to a validated problem/JTBD rather than simply inherited from the BRD.