

Tattava Copilot — V1 Product Definition & Planning

Source set based on

1. **Business Requirements Document (BRD) v1.0**
 2. **Product Discovery Artifact v1.0**
 3. **AI-Native Product Strategy v1.0**
-

Executive Summary

1. Product Definition

Tattava Copilot V1 is an **AI-native Story Development and Project Intelligence system** for entertainment teams.

Its core job is not simply to generate scripts.

It should maintain a persistent representation of a project — its:

- premise
- story
- characters
- relationships
- research
- decisions
- canon
- narrative structure
- versions
- approvals

—and use that project state to help humans research, reason, develop, write, evaluate and maintain continuity.

This follows the strategic distinction that Tattava should understand project information and reason over it rather than merely store it or answer isolated prompts.

2. V1 North-Star Workflow

```
graph TD; A[CREATIVE INPUT] --> B[PROJECT CREATION]; B --> C[INTENT + CONTEXT]; C --> D[STORY BRAIN]; D --> E[RESEARCH + EVIDENCE]; E --> F[CHARACTER INTELLIGENCE]; F --> G[NARRATIVE REASONING]; G --> H[STORY DEVELOPMENT]; H --> I[FORMAT / GENRE];
```

↓
WRITING
↓
SCREENPLAY
↓
DIALOGUE
↓
CANON / CONTINUITY
↓
EVALUATION
↓
HUMAN APPROVAL
↓
STRUCTURED STORY DEVELOPMENT PACKAGE

3. Recommended Product Architecture

Recommendation: build V1 around one central object:

Project Intelligence / Story Brain

rather than around 14 disconnected agents.



The strategic principle is explicit: **AI performs cognitive work; humans retain creative authority.**

1. V1 Product Blueprint

1.1 Product Vision

Source-backed vision

Tattava's long-term vision is an AI-native operating system for entertainment development and production that transforms ideas into structured, production-ready projects while preserving human creative control.

V1 interpretation

Tattava V1 is the first intelligence layer of that operating system.

It should enable a creative team to take an initial idea and progressively transform it into a:

| **researched → structured → context-aware → continuity-aware → human-approved Story Development Package.**

1.2 V1 Objective

Primary objective

| **Reduce the time and operational complexity required to move an entertainment idea toward a human-approved structured story package while preserving human creative authority.**

For V1, the practical measurement proxy is:

| **Time from accepted project input → human-approved Story Development Package.**

This is directly aligned with Discovery's proposed V1 proxy.

1.3 Target Users

Priority	Persona	Primary Job
P0	Story Developer / Story Editor	Develop coherent narrative
P0	Writer / Screenwriter	Develop and transform story into written material
P0	Creative Producer	Coordinate development and progression
P0	Director / Showrunner	Provide creative direction and approve decisions
P1	Producer	Track project progression
P2	Production / Art / Specialist teams	Downstream project intelligence

The Discovery artifact identifies Writers, Story Developers/Editors, Creative Producers and Directors/Showrunners as the principal P0 discovery users.

1.4 Core Problems

The product should address five connected problems:

Problem 1 — Project-context fragmentation

Information exists across:

- documents
- notes
- research

- scripts
- conversations
- people
- AI chats

Users repeatedly reconstruct context.

Problem 2 — Repetitive development work

Teams repeatedly:

- research
- summarize
- restructure
- rewrite
- convert
- compare
- consolidate
- document

Problem 3 — Research traceability

Research may be disconnected from the creative decisions it informs.

Problem 4 — Continuity risk

Characters, facts, timelines and story logic can drift between iterations.

Problem 5 — Creative handoff friction

Research → Story → Writing → Screenplay → Dialogue creates repeated translation and context loss.

These problems are explicitly identified in Discovery.

1.5 Core JTBD

JTBD	Desired Outcome
Develop an idea	Premise, conflict, characters, themes, possibilities
Conduct meaningful research	Relevant, credible, traceable evidence
Maintain project context	Persistent project memory
Develop characters	Consistent motivation, relationships and voice
Develop story structure	Better-informed narrative decisions
Write with context	Less generic/inconsistent AI output
Review and diagnose	Identify structural and continuity issues
Prepare development package	Consolidated, review-ready project

These eight JTBDs are explicitly defined in the Discovery Artifact.

1.6 Value Proposition

For creative professionals

Stop repeatedly explaining the project to the AI. Tattava remembers the project, understands its current state, connects research to decisions, helps develop the story and flags contradictions before they propagate.

For creative producers

Move projects forward with a structured, traceable project state rather than fragmented documents and conversations.

For directors/showrunners

Explore alternatives and review consequences while retaining final creative authority.

1.7 Core V1 Capabilities

The BRD defines 14 V1 capabilities.

#	Capability	Product role	Recommended priority
1	Input Layer	Establish project context	P0
2	Story Brain	Persistent project intelligence	P0
3	Character Intelligence	Structured character context	P0
4	Narrative Reasoning	Story logic analysis	P0
5	Story Development	Narrative generation/development	P0
6	Research	Evidence + source intelligence	P0
7	Format / Genre	Contextual format decisions	P1
8	Template Selection	Structure downstream work	P1
9	Writing Interface	Human + AI writing workflow	P0
10	Screenplay Agent	Convert approved story into screenplay structure	P0/P1
11	Dialogue Agent	Context-aware dialogue	P1
12	Story Doctor	Critique and repair	P1
13	Canon / Continuity	Detect contradictions	P0
14	Evaluation	Assess readiness	P0

Why not simply make all 14 P0?

Because Discovery explicitly states that the exact MVP has **not been validated** and that the BRD's complete capability list should not automatically become engineering scope.

1.8 What AI Should Do

AI should:

- understand user intent
- identify missing information
- retrieve relevant project context
- conduct/synthesize research
- generate alternatives
- identify narrative issues
- construct character intelligence
- draft/rewrite material

- transform approved content
- identify continuity conflicts
- explain evidence
- evaluate against defined criteria
- propose next actions
- surface uncertainty

1.9 What AI Should NOT Do

AI should not independently:

- approve creative direction
- alter approved canon silently
- overwrite approved material
- make irreversible creative decisions
- promote unverified research into authoritative project facts
- access unauthorized project information
- make production/business decisions outside V1
- silently modify project state

This follows the BRD's human-in-the-loop principle.

1.10 Human-in-the-Loop Model

The fundamental operating loop should be:

```

AI proposes
↓
Human reviews
↓
Human edits / accepts / rejects
↓
System records decision
↓
Approved state becomes project intelligence
↓
Future AI uses approved state

```


This is much more important than maximizing AI autonomy.

The Discovery hypothesis explicitly describes the preferred pattern as:

AI proposes → Human evaluates → Human modifies → Human approves → Project updates.

1.11 Conceptual Architecture





The Strategy explicitly recommends that durable value sit **above the models**, in project intelligence, workflow and orchestration, rather than in the underlying model.

1.12 Success Metrics

North-star proxy

Time to Human-Approved Story Development Package

Supporting metrics

Efficiency

- Research turnaround
- Time to first useful artifact
- Time per development iteration

- Time spent reconstructing context
- Repetitive work saved

Quality

- Human acceptance rate
- Correction rate
- Canon consistency
- Continuity issues detected
- Missed contradiction rate
- Research verification rate

Trust

- AI rejection rate
- Verification rate
- Approval rate
- Source interaction

Adoption

- Projects created
- Returning projects
- Active projects
- Projects reaching Story Brain activation
- Projects reaching approved output

These metric categories are supported by Discovery and Strategy.



2. PRD

2.1 Problem Statement

Entertainment development requires repeated transformation of information across research, story, characters, writing, screenplay, dialogue and review.

The problem is not merely content generation.

The deeper problem is that project intelligence can become fragmented, causing:

- repeated context-setting
- repetitive work
- disconnected research
- version confusion
- continuity risk
- collaboration friction

- reduced traceability.

This is the central problem formulation in Discovery.

2.2 Product Goals

G1 — Accelerate development

Reduce time between initial idea and structured story output.

G2 — Establish persistent project intelligence

Make project context reusable across the workflow.

G3 — Improve research trust

Connect evidence to sources and project relevance.

G4 — Improve continuity

Maintain canon and detect contradictions.

G5 — Preserve creative control

Make humans the authority over important creative decisions.

G6 — Establish foundation for downstream versions

Design the intelligence model so later visual, production and localization workflows can consume it.

2.3 Non-Goals

V1 should **not** attempt to build:

- AI video generation
- AI filmmaking
- storyboarding/visual production
- localization
- dubbing
- production scheduling
- casting management
- budgeting
- production tracking
- rights/revenue management
- complete production OS
- standalone image/video generator

The BRD roadmap explicitly places these capabilities beyond V1.

2.4 Epics

Epic	Objective
EPIC-01 Project Creation	Establish project

Epic	Objective
EPIC-02 Project Intelligence	Build Story Brain
EPIC-03 Research Intelligence	Research + provenance
EPIC-04 Character Intelligence	Maintain character state
EPIC-05 Narrative Intelligence	Reason over story
EPIC-06 Story Development	Develop story
EPIC-07 Writing Workspace	Human + AI writing
EPIC-08 Screenplay	Transform approved story
EPIC-09 Canon & Continuity	Protect consistency
EPIC-10 Evaluation	Diagnose readiness
EPIC-11 Approval & Versioning	Govern state
EPIC-12 Final Package	Produce structured output

2.5 Functional Requirements

EPIC-01 — Project Creation

FR-PROJ-001

System shall allow users to create a project.

FR-PROJ-002

System shall accept:

- idea
- logline
- concept
- script/reference material
- character ideas
- themes
- creative brief

These input types are explicitly identified in the BRD.

FR-PROJ-003

System shall establish initial project context from supplied material.

FR-PROJ-004

System shall identify missing or ambiguous context.

EPIC-02 — Story Brain

FR-BRAIN-001

System shall maintain persistent project intelligence.

FR-BRAIN-002

Story Brain shall support at minimum:

- project
- premise
- story
- characters
- relationships
- research
- decisions
- canon
- versions
- approvals

FR-BRAIN-003

Downstream AI operations shall retrieve relevant project context automatically.

FR-BRAIN-004

Approved project state shall be distinguishable from AI-generated suggestions.

EPIC-03 — Research

FR-RESEARCH-001

System shall identify research questions relevant to the project.

FR-RESEARCH-002

System shall retrieve/synthesize relevant research.

FR-RESEARCH-003

Every research finding should support:

```
Source
Reference
Finding
Project relevance
```

This provenance structure is explicitly defined in the BRD.

FR-RESEARCH-004

System shall distinguish:

- sourced evidence
- AI inference
- unresolved uncertainty.

FR-RESEARCH-005

Research findings may be promoted into project intelligence only through defined state rules.

EPIC-04 — Character Intelligence

FR-CHAR-001

System shall identify characters from project material.

FR-CHAR-002

System shall maintain:

- motivation
- voice hypotheses
- relationships
- relevant traits

FR-CHAR-003

Character information shall be available to downstream narrative and writing operations.

EPIC-05 — Narrative Reasoning

FR-NARR-001

System shall reason over:

- cause/effect
- motivation
- structure
- narrative logic

FR-NARR-002

System shall flag:

- plot holes
- motivation gaps
- structural weaknesses.

FR-NARR-003

AI findings must be presented as recommendations/diagnostics, not authoritative creative decisions.

EPIC-06 — Story Development

FR-STORY-001

System shall assist with:

- premise
- logline
- synopsis
- story structure
- themes
- conflict
- relationships

- narrative direction
- beats.

FR-STORY-002

User shall be able to:

- accept
- reject
- modify
- regenerate.

FR-STORY-003

Approved output shall update project intelligence.

EPIC-07 — Writing

FR-WRITE-001

Writer interface shall support:

- review
- editing
- rewriting
- expansion
- shortening
- refinement
- regeneration
- creative direction.

This is explicitly required by the BRD.

FR-WRITE-002

AI generation must be grounded in current approved project context.

FR-WRITE-003

AI shall not silently overwrite approved content.

EPIC-08 — Screenplay

FR-SCRIPT-001

System shall transform sufficiently developed/approved story material into screenplay-oriented structure.

FR-SCRIPT-002

Supported outputs may include:

- scene structure
- scene descriptions
- transitions

- action descriptions
- character appearances
- dialogue placeholders.

These are BRD-defined screenplay capabilities.

EPIC-09 — Canon & Continuity

FR-CANON-001

System shall maintain established facts.

FR-CANON-002

Canon shall cover:

- characters
- world
- timeline
- story logic.

FR-CANON-003

System shall flag contradictions.

FR-CANON-004

Human shall determine whether a contradiction is actually a problem.

The BRD defines Canon & Continuity as maintaining established facts and checking new material against them.

EPIC-10 — Evaluation

FR-EVAL-001

System shall evaluate story outputs against defined criteria.

Potential dimensions:

- structure
- genre fit
- format fit
- continuity
- narrative logic
- completeness.

FR-EVAL-002

Evaluation shall produce:

- strengths
- gaps
- risks
- recommended improvements.

FR-EVAL-003

Evaluation shall not automatically approve a project.

EPIC-11 — Approval / Versioning

FR-APP-001

Users shall approve/reject AI-generated outputs.

FR-APP-002

Material changes should create versions rather than silently replacing approved state.

FR-APP-003

System shall preserve approval history.

FR-APP-004

System shall distinguish:

```
AI proposal
Human edit
Human approval
Approved project state
```

EPIC-12 — Final Package

The final package should support the BRD-defined structure:

- project concept
- research
- logline
- synopsis
- genre/format
- story structure
- character overview
- screenplay
- dialogue
- supporting research
- evaluation
- project metadata / Story Brain state.

However:

Recommendation: do not require every artifact to be complete before a project can be considered useful.

2.6 AI Requirements

For every AI capability, Tattava should maintain this contract:

Dimension	Requirement
Intent	Understand what user wants
Context	Retrieve relevant project state
Input	User request + project context
Reasoning	Determine relevant evidence and constraints
Tools	Search/retrieval/project actions where authorized
Output	Structured candidate result
Confidence	Communicate uncertainty
Approval	Human confirmation where state changes
Failure	Fallback / retry / clarification
Observability	Record AI run, context and outcome

This directly implements the Discovery AI-native test: **intelligence** → **context** → **reasoning** → **action** → **memory** → **feedback** → **evaluation** → **control**.

2.7 User Stories

US-001 — Create project

As a writer, I want to create a project from an idea so that I can begin structured development.

Acceptance criteria

- Project is created.
- Initial input is stored.
- Story Brain is initialized.
- Missing information is identified.
- User can edit the initial interpretation.

US-002 — Ask Tattava to research

As a story developer, I want Tattava to research a project-specific question so that I can make an informed creative decision.

Acceptance criteria

- Research question is identified.
- Relevant sources are retrieved.
- Findings contain provenance.
- Project relevance is shown.
- Unsupported claims are not presented as verified facts.

US-003 — Develop story

As a story developer, I want Tattava to generate story alternatives using my project's established context.

Acceptance criteria

- Existing project context is used.
- Alternatives are clearly separated.

- User can modify/reject/approve.
- Approved result updates project intelligence.

US-004 — Write with context

| As a writer, I want AI writing grounded in approved project context so that I spend less time correcting generic output.

Acceptance criteria

- Relevant character/story/canon information is retrieved.
- Generated content reflects approved state.
- User can edit/reject/regenerate.
- Approved output becomes versioned project content.

US-005 — Detect continuity issue

| As a story editor, I want Tattava to flag contradictions so that I can review them before they propagate.

Acceptance criteria

- Contradiction includes affected entities.
- System shows conflicting information.
- User can inspect source versions.
- User can dismiss or resolve.
- No automatic canon change occurs.

2.8 UX Requirements

Tattava should not primarily look like:

```
ChatGPT
+
Sidebar
+
Generate button
```

Recommended UX model

```
PROJECT
|
├─ Overview
├─ Story Brain
├─ Research
├─ Characters
├─ Story
├─ Writing
├─ Screenplay
├─ Canon
├─ Evaluation
├─ Versions
└─ Final Package
```

The AI should appear **inside workflows**, not as a separate destination.

Example:

Character: Arjun

Motivation

[Current approved state]

Relationships

[Current project state]

Canon

[Established facts]


AI Suggestions

- Possible motivation conflict
- Research relevant to character
- 3 development directions

[Accept] [Edit] [Reject]

2.9 Data Requirements

Minimum conceptual entities:



Critical distinction

Do not treat the Story Brain as a giant prompt or chat transcript.

It should be a **structured project state**.

This follows Strategy's explicit definition of Project Intelligence as entities, relationships, decisions, evidence, versions, dependencies, approvals and creative intent.

2.10 Permissions

At minimum:

Role	View	Edit	Generate	Approve
Writer	✓	✓	✓	TBD
Story Developer	✓	✓	✓	TBD
Story Editor	✓	✓	✓	TBD
Director	✓	✓	✓	✓
Creative Producer	✓	✓	✓	TBD

Role	View	Edit	Generate	Approve
Admin	✓	✓	✓	Policy-defined

Not specified in source artifacts: exact organizational approval hierarchy by stage.

Therefore this must remain an **Open Question**, not an invented requirement.

2.11 Security / Privacy / IP

Because film projects contain sensitive unreleased IP, V1 should include:

- project isolation
- authorized access
- secure storage
- controlled AI context retrieval
- auditability
- version history
- controlled external research/tool access
- no unauthorized cross-project retrieval.

The BRD explicitly identifies security, IP and access control as requirements/risk areas.

2.12 Edge Cases

Empty/weak input

```
User gives:
"My story is about a village."

AI:
"I can start, but several important decisions are unspecified."

→ Ask targeted questions
→ Offer optional assumptions
→ Never silently invent them as canon
```

Conflicting project information

```
Character Age = 32
Later draft = 35

AI:
"Potential continuity conflict detected."

[Keep 32]
[Change to 35]
[Create new version]
[Ignore]
```

Research uncertainty

```
Evidence insufficient
↓
AI explicitly reports uncertainty
↓
Suggests additional research
```

↓
Does not promote claim to verified project fact

AI failure

Research unavailable
↓
Show failure
↓
Preserve existing project state
↓
Allow retry / alternative action

Context overload

AI should not blindly retrieve the entire project.

It should retrieve **relevant project state**.

3. MVP Scope

This is the most important section for preventing scope creep.

3.1 MVP Definition

Recommendation:

MVP should prove one hypothesis:

A persistent, project-grounded AI workflow can help a creative professional move from idea → research → story development → context-aware writing → continuity review → human-approved package more effectively than fragmented tools.

This is directly aligned with the Discovery MVP hypothesis.

3.2 MVP Must-Haves

Capability	MVP
Project creation	✓
Initial input	✓
Story Brain	✓
Project context retrieval	✓
Research + sources	✓
Character intelligence	✓
Story development	✓
Context-aware writing	✓
Basic canon/continuity	✓
Human approval	✓
Versioning	✓
Structured final package	✓
Basic evaluation	✓

3.3 MVP Should-Haves

- Intent understanding
 - Format/genre recommendation
 - Basic screenplay transformation
 - Basic dialogue generation
 - Basic Story Doctor
 - Decision log
 - Better dependency intelligence
 - Project-level analytics
-

3.4 Post-MVP

- Advanced dependency graph
 - Advanced collaboration
 - Full Story Doctor
 - Deep screenplay intelligence
 - Advanced dialogue intelligence
 - Story World
 - Visual development
 - Storyboarding
 - Image generation
 - Video generation
 - Production planning
 - Scheduling
 - Budgeting
 - Casting
 - Localization
 - Distribution
 - Business/management OS
-

3.5 Explicitly Out of Scope

```
graph TD; V1 --> SD[Story Development]; SD --> R[Research]; R --> W[Writing]; W --> S[Screenplay]; S --> C[Continuity]; C --> P[Package];
```

```
X
├─ Video generation
├─ Image generation platform
├─ Production scheduling
├─ Budgeting
├─ Casting
├─ Localization
├─ Distribution
└─ Full production OS
```

The BRD explicitly places visual development, AI filmmaking, localization, production operations and MVV OS functionality beyond V1.

3.6 MVP Feature Decision Logic

Story Brain

User value: eliminates repeated context setting.

Business value: establishes Tattava's strategic differentiator.

Why MVP: foundational to almost every other capability.

Risk: incorrect memory could contaminate downstream outputs.

Research

User value: relevant evidence with provenance.

Business value: establishes trust.

Why MVP: one of the strongest differentiated hypotheses.

Risk: hallucination / source quality.

Story Development

User value: faster exploration.

Business value: directly moves project toward output.

Why MVP: core user job.

Risk: generic AI output.

Writing

User value: context-aware transformations.

Business value: demonstrates intelligence beyond research.

Why MVP: bridges project intelligence to actual creative work.

Risk: writer rejects output because it does not preserve intent.

Canon / Continuity

User value: catches contradictions.

Business value: validates persistent intelligence.

Why MVP: differentiates project intelligence from ordinary chat.

Risk: false positives.

Human Approval


User value: creative control.

Business value: trust/adoption.

Why MVP: foundational product principle.


3.7 MVP Golden Loop

The MVP should optimize this loop:



4. User Flows

Flow 1 — Create Project



AI behavior

AI:


- extracts intent

- identifies entities
- detects ambiguity
- proposes initial structure

AI does **not**:

- decide creative direction
- silently establish disputed facts as canon.

Flow 2 — Research



Success state

Research is no longer merely a document.


It becomes:

```

Source
↓
Evidence
↓
Finding
↓
Project relevance
↓
Story implication
↓
Decision
  
```

This research-to-story intelligence model is explicitly proposed in Strategy.


Flow 3 — Story Development




Important

The AI's output is a **candidate** until approved.

Flow 4 — Context-Aware Writing




Flow 5 — Continuity Check



Key design principle

Continuity should operate throughout the workflow, not only as a final "check my script" button. This is explicitly stated in the Strategy Artifact.


Flow 6 — Final Package




5. AI-Native Product Architecture & Behavior

5.1 The AI-Native Loop

The product should behave as:



Not:



5.2 Deterministic vs AI Behavior

Function	Deterministic	AI
User authentication	✓	
Permission checks	✓	
Project isolation	✓	
Version creation	✓	
Approval state	✓	
Audit log	✓	
Package assembly	✓	
Context retrieval rules	✓	✓
Intent understanding		✓
Research synthesis		✓
Story generation		✓
Character reasoning		✓
Narrative reasoning		✓
Continuity detection		✓
Evaluation		✓
Creative approval		Human

5.3 AI Capability Contracts

Research Agent

Intent: obtain reliable evidence.

Context: project + research question.

Inputs: query, project context.

Tools: search/retrieval.

Output: evidence-backed findings.

Boundary: cannot treat unsupported claims as facts.

Human: verifies important findings.

Failure: insufficient evidence.

Evaluation: citation correctness, unsupported claim rate.

Story Agent

Intent: develop narrative.

Context: Story Brain + research + characters + canon.

Output: alternatives / story structures / beats.

Boundary: does not approve creative direction.

Human: chooses/edits.

Evaluation: acceptance, edits, regeneration rate.

Character Intelligence

Intent: maintain coherent character model.

Context: all approved character-related project information.

Output: structured character state.

Failure: conflicting or incomplete evidence.

Human: resolves ambiguity.

Continuity Agent

Intent: identify contradictions.

Context: canon + current artifact + versions.

Output:

```
Conflict
Evidence A
Evidence B
Affected entity
Potential impact
Confidence
```

Boundary: does not automatically rewrite canon.

Evaluation:

- precision
 - recall
 - false positives
 - missed issues
 - review time.
-

Evaluation Agent

Intent: diagnose readiness.

Output:

- strengths
- weaknesses
- risks
- missing elements
- suggested improvements

Boundary: "evaluation" is not "approval."

5.4 Trust Model

Every important AI output should answer:

Where did this come from?

Source / project context.

Why is this being suggested?

Reason / relevant context.

How certain is it?

Confidence or uncertainty.

What does it affect?

Affected project entities.

Who approved it?

Human approval record.

When did it change?

Version history.

This follows Strategy's traceability principle that important artifacts should expose where information came from, when it changed, why it changed, who approved it and what it affects.

6. Cross-Artifact Traceability Matrix

Requirement / Capability	BRD	Discovery	Strategy	V1	PRD	MVP	User Flow
Project creation	✓	✓	✓	✓	✓	✓	✓
Persistent Story Brain	✓	✓	✓	✓	✓	✓	✓
Research	✓	✓	✓	✓	✓	✓	✓
Research provenance	✓	✓	✓	✓	✓	✓	✓
Character intelligence	✓	✓	✓	✓	✓	✓	✓

Requirement / Capability	BRD	Discovery	Strategy	V1	PRD	MVP	User Flow
Narrative reasoning	✓	✓	✓	✓	✓	Partial	✓
Story development	✓	✓	✓	✓	✓	✓	✓
Format / genre	✓	✓	✓	✓	✓	Should	—
Templates	✓	✓	—	✓	✓	Should	—
Context-aware writing	✓	✓	✓	✓	✓	✓	✓
Screenplay	✓	✓	✓	✓	✓	Should/P0 decision	✓
Dialogue	✓	✓	✓	✓	✓	Should	—
Story Doctor	✓	✓	✓	✓	✓	Should	—
Canon / continuity	✓	✓	✓	✓	✓	✓	✓
Evaluation	✓	✓	✓	✓	✓	✓	✓
Human approval	✓	✓	✓	✓	✓	✓	✓
Versioning	✓	✓	✓	✓	✓	✓	✓
Final package	✓	✓	✓	✓	✓	✓	✓
Visual development	Future	Future	Future	No	No	No	No
AI video	Future	Future	Future	No	No	No	No
Production planning	Future	Future	Future	No	No	No	No
Localization	Future	Future	Future	No	No	No	No

6.1 Major Cross-Artifact Conflict #1 — What exactly is V1?

BRD

Defines V1 as a broad 14-capability pipeline.

Discovery

States:

| exact V1 MVP is not validated.

and proposes:

| Project Intelligence + Research + Context-Aware Story Development + Canon/Continuity + Human Approval.

Strategy

Defines the Golden Workflow as:

| Idea → Story Brain → Research → Story Development → Writing → Screenplay → Continuity → Evaluation → Approved Package.

Resolution

Do not collapse V1 and MVP into one concept.

Use:

```
V1 PRODUCT DEFINITION
↓
Broad capability architecture
↓
MVP RELEASE
↓
Validated P0 golden workflow
↓
P1 capabilities
```

This preserves the BRD while respecting Discovery evidence.

6.2 Major Conflict #2 — Screenplay / Dialogue priority

BRD treats Screenplay and Dialogue as V1 capabilities.

Discovery includes them in the proposed workflow but explicitly says the exact smallest useful V1 is unvalidated.

Strategy's Phase 1 focus emphasizes:

- Project Creation
- Story Brain
- Research
- Character Intelligence
- Narrative Reasoning
- Story Development
- Writer Interface
- Canon
- Evaluation

while its Golden Workflow includes Screenplay and Dialogue.

Minimum decision required

Management/Product decision:

Is screenplay transformation required for the first pilot, or is context-aware story/writing + continuity sufficient to prove the core hypothesis?

Do not allow engineering scope to decide this implicitly.

6.3 Major Conflict #3 — Current workflow evidence

Discovery repeatedly states that the actual workflow, tools, pain severity, AI usage and switching behavior are unknown.

Yet Strategy presents a current-state workflow:

```
Idea
↓
Discussion
↓
Google/Web Research
↓
Notes/Docs
↓
Story
↓
```

Characters
↓
Writer drafts
↓
Feedback
↓
Revision
↓
Screenplay
↓
Continuity

The Strategy correctly describes this as consistent with the BRD, but Discovery says the current workflow remains largely inferred.

Treatment

Inference: Strategy's current-state workflow should be used as a **working hypothesis**, not as validated user behavior.

7. Risks, Assumptions & Open Questions

7.1 High-Risk Assumptions

Assumption	Status	Validation
Context fragmentation is significant	Hypothesis	User interviews
Story Brain creates meaningful value	Hypothesis	Prototype
Research traceability increases trust	Hypothesis	Research experiment
Continuity saves meaningful time	Hypothesis	Task test
Integrated workflow beats separate tools	Hypothesis	End-to-end prototype
Users will change workflow	Unknown	Switching interviews
V1 can create standalone value	Medium-confidence hypothesis	V1 prototype

These are directly reflected in Discovery's assumption map.

7.2 AI Risks

Hallucinated research

Mitigation: sources + evidence + provenance.

Generic creative output

Mitigation: project-grounded generation.

Incorrect Story Brain

Mitigation: human approval + explicit state distinction.

Continuity false positives

Mitigation: show evidence and let human decide.

Context contamination

Mitigation: structured retrieval rather than dumping full history.

AI overreach

Mitigation: deterministic approval gates.

7.3 Product Risks

Scope explosion

14 capabilities can easily become 14 products.

Mitigation: Golden Workflow + MVP gate.

Chatbot trap

Tattava could become "ChatGPT with project files."

Mitigation: Project Intelligence as the primary product object.

AI-first instead of workflow-first

Building agents without solving the workflow.

Mitigation: every AI capability must map to a JTBD and measurable outcome.

Premature architecture

Building full V4 production infrastructure during V1.

Mitigation: future-compatible data model, but V1 delivery scope remains narrow.

7.4 Open Questions

User / Workflow

1. Who actually owns each creative decision?
2. Who approves Story Brain changes?
3. Where are current drafts stored?
4. What tools are currently used?
5. How are revisions managed?
6. What is the largest actual time sink?
7. How often is AI currently used?
8. Which AI tasks are trusted?
9. Which AI tasks are rejected?
10. What causes users to switch workflows?

Discovery explicitly identifies these as unresolved questions.

Product

1. What is the exact minimum package required for a project to be considered "developed"?
2. Is screenplay generation P0?
3. Is dialogue generation P0?
4. Is Story Doctor P0?
5. How much autonomy should each agent have?
6. What exactly qualifies as "approved canon"?
7. What is the approval hierarchy?
8. What constitutes a production-ready story package?

Technical

1. Which LLM/model providers will be used?
2. What external research sources are permitted?
3. What is the retention policy?
4. What data can be sent to external model providers?
5. What is the exact permission model?
6. What is the project/version schema?
7. What is the evaluation benchmark for AI outputs?

Not specified in source artifacts: answers to these technical questions.

8. Recommended Decisions Before Development

Decision 1 — Separate V1 from MVP

Recommendation

Use:

▮ **V1 = product architecture/capability definition**

and

▮ **MVP = smallest validated release.**

This is the cleanest resolution of the BRD/Discovery tension.

Decision 2 — Make Story Brain the central product object

Do not build:

```
Research Agent
Story Agent
Writing Agent
Continuity Agent
```

as isolated products.

Build:

```
PROJECT
↓
STORY BRAIN
↓
Agents reason over project state
↓
Human approval
↓
Project state evolves
```

This is the central strategic thesis.

Decision 3 — Define the MVP around one Golden Workflow

Recommended MVP:

```
Idea
↓
Project
↓
Story Brain
↓
Research
↓
Characters
↓
Story Development
↓
Context-Aware Writing
↓
Canon / Continuity
↓
Evaluation
↓
Human Approval
↓
Story Package
```

Screenplay and Dialogue should be included only to the depth required to validate the workflow, unless management explicitly makes them P0.

Decision 4 — Establish approval semantics before engineering

Define:

```
AI Generated
↓
Candidate
↓
Human Edited
↓
Human Approved
↓
Canonical Project State
```

Without this, the Story Brain can become contaminated by unapproved AI content.

Decision 5 — Define the Project Intelligence schema before building agents

The schema should be designed around:

```
Project
Story
Character
Relationship
Research
Source
Decision
Canon
Artifact
Version
Approval
Evaluation
Dependency
```

The AI layer should consume this structured state.

Decision 6 — Establish AI evaluation before launch

Do not wait until after the product is built.

Create benchmark datasets for:

Research

- source correctness
- citation correctness
- unsupported claims

Story

- project grounding
- user acceptance
- edit distance

Character

- consistency

Continuity

- contradiction precision
- contradiction recall
- false positives

Overall

- time-to-usable-output
- time-to-approved-package.

Discovery already provides the foundations for these experiments.

Final Product Definition

The three artifacts converge on a clear strategic direction:

Tattava should not be built as 14 independent AI features. It should be built as one intelligent project-development system whose capabilities reinforce a persistent Project Intelligence Layer.

The core V1 equation is:

```
PROJECT
+
PROJECT INTELLIGENCE
+
RESEARCH
+
AI REASONING
+
CONTEXT-AWARE CREATION
+
CANON / CONTINUITY
+
HUMAN APPROVAL
↓
STRUCTURED STORY DEVELOPMENT PACKAGE
```

The strongest differentiation is therefore **not "better AI writing."**

It is:

```
Generic AI
Prompt
↓
Response
```

versus:

```
TATTVA
Project
↓
Persistent Intelligence
↓
Evidence
↓
Creative Decisions
↓
Story State
↓
AI Reasoning
↓
Generation
↓
Validation
↓
Human Approval
↓
Updated Project Intelligence
```

That is the AI-native product architecture supported across the BRD, Discovery and Strategy artifacts. The Discovery document's final thesis is particularly important: the central hypothesis remains something to **prove**, not a validated market claim.

The immediate product-development sequence should therefore be:

```
SOURCE ARTIFACTS
↓
V1 PRODUCT DEFINITION    ← THIS DOCUMENT
↓
MVP / P0 DECISION
↓
UX / PROTOTYPE
↓
USER VALIDATION
↓
PRD FREEZE
↓
AI EVALUATION SPEC
↓
ARCHITECTURE BLUEPRINT
↓
ENGINEERING BUILD SPEC
↓
IMPLEMENTATION
↓
PILOT
↓
MEASURE
↓
ITERATE
```

The single most important unresolved decision before engineering starts is the exact P0 boundary of Screenplay, Dialogue, Story Doctor, Format/Genre and Evaluation. The source artifacts agree on the strategic destination but intentionally leave the smallest validated release open. That should be resolved through the Discovery → Prototype → Validation gates rather than by allowing the entire BRD capability list to become MVP scope.